

CIND 123: Data Analytics Basic Methods: Assignment-3

Assignment 3 (10%)

Total 100 Marks

[Stephanie Boissonneault]

Instructions

This is an R Markdown document. Markdown is a simple formatting syntax for authoring HTML, PDF, and MS Word documents. For more details on using R Markdown see <http://rmarkdown.rstudio.com>.

Use RStudio for this assignment. Complete the assignment by inserting your R code wherever you see the string “#INSERT YOUR ANSWER HERE”.

When you click the **Knit** button a document will be generated that includes both content as well as the output of any embedded R code chunks within the document. You can embed an R code chunk like this:

Submit **both** the rmd and generated output files. Failing to submit both files will be subject to mark deduction.

Sample Question and Solution

Use `seq()` to create the vector $(2, 4, 6, \dots, 20)$.

```
#INSERT YOUR ANSWER HERE.  
seq(2,20,by = 2)
```

```
## [1]  2  4  6  8 10 12 14 16 18 20
```

Question 1 [15 Pts]

- a) [5 Pts] First and second midterm grades of some students are given as `c(85,76,78,88,90,95,42,31,66)` and `c(55,66,48,58,80,75,32,22,39)`. Set R variables `first` and `second` respectively. Then find the least-squares line relating the second midterm to the first midterm.

Does the assumption of a linear relationship appear to be reasonable in this case? Give reasons to your answer as a comment.

```
first <- c(85,76,78,88,90,95,42,31,66)  
second <- c(55,66,48,58,80,75,32,22,39)  
  
lsmodel <- lm(second ~ first)  
summary(lsmodel)
```

```
##
## Call:
## lm(formula = second ~ first)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9.238 -7.747  1.753  4.383 13.318
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -4.1516    10.9987  -0.377  0.71702
## first         0.7870     0.1461   5.389  0.00102 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 9.175 on 7 degrees of freedom
## Multiple R-squared:  0.8058, Adjusted R-squared:  0.778
## F-statistic: 29.04 on 1 and 7 DF,  p-value: 0.001021
```

#The assumption of a linear relationship appears to be reasonable because the summary of the line of be


- b) [5 Pts] Plot the second midterm as a function of the first midterm using a scatterplot and graph the least-square line in red color on the same plot.

```
print(plot(first, second, main = "First and second midterm grades of some students"))
```

```
## NULL
```

```
print(abline(lsmode1, col = "red"))
```

First and second midterm grades of some students



```
## NULL
```

- c) [5 Pts] Use the regression line to predict the second midterm grades when the first midterm grades are 81 and 23.

```
first.81 <- data.frame(first=c(81))
predict(lsmode, first.81)
```

```
##          1
## 59.59881
```

```
first.23 <- data.frame(first=c(23))
predict(lsmode, first.23)
```

```
##          1
## 13.95039
```

Question 2 [45 Pts]

This question makes use of package “plm”. Please load Crime dataset as follows:

```
#install.packages("plm")
library(plm)
```

```
## Warning: package 'plm' was built under R version 4.3.2
```

```
data(Crime)
```

- a) [5 Pts] Display the first 8 rows of 'crime' data and display the names of all the variables, the number of variables, then display a descriptive summary of each variable.

```
#Display first 8 rows
print(head(Crime, 8))
```

```
##      county year      crmrte  prbarr  prbconv  prbpris avgsen      polpc  density
## 1         1    81 0.0398849 0.289696 0.402062 0.472222   5.61 0.0017868 2.307159
## 2         1    82 0.0383449 0.338111 0.433005 0.506993   5.59 0.0017666 2.330254
## 3         1    83 0.0303048 0.330449 0.525703 0.479705   5.80 0.0018358 2.341801
## 4         1    84 0.0347259 0.362525 0.604706 0.520104   6.89 0.0018859 2.346420
## 5         1    85 0.0365730 0.325395 0.578723 0.497059   6.55 0.0019244 2.364896
## 6         1    86 0.0347524 0.326062 0.512324 0.439863   6.90 0.0018952 2.385681
## 7         1    87 0.0356036 0.298270 0.527596 0.436170   6.71 0.0018279 2.422633
## 8         3    81 0.0163921 0.202899 0.869048 0.465753   8.45 0.0005939 0.976834
##      taxpc  region smsa  pctmin      wcon      wtuc      wtrd      wfir      wser
## 1 25.69763 central   no 20.21870 206.4803 333.6209 182.3330 272.4492 215.7335
## 2 24.87425 central   no 20.21870 212.7542 369.2964 189.5414 300.8788 231.5767
## 3 26.45144 central   no 20.21870 219.7802 1394.8030 196.6395 309.9696 240.1568
## 4 26.84235 central   no 20.21870 223.4238 398.8604 200.5629 350.0863 252.4477
## 5 28.14034 central   no 20.21870 243.7562 358.7830 206.8827 383.0707 261.0861
## 6 29.74098 central   no 20.21870 257.9139 369.5465 218.5165 409.8842 269.6129
## 7 30.99368 central   no 20.21870 281.4259 408.7245 221.2701 453.1722 274.1775
## 8 14.56088 central   no 7.91632 188.7683 292.6422 151.4234 202.4292 191.3742
##      wmfgr  wfed  wsta  wloc      mix  pctymle  lcrmrte  lprbarr
## 1 229.12 409.37 236.24 231.47 0.0999179 0.0876968 -3.221757 -1.238923
## 2 240.33 419.70 253.88 236.79 0.1030491 0.0863767 -3.261134 -1.084381
## 3 269.70 438.85 250.36 248.58 0.0806787 0.0850909 -3.496449 -1.107303
## 4 281.74 459.17 261.93 264.38 0.0785035 0.0838333 -3.360270 -1.014662
## 5 298.88 490.43 281.44 288.58 0.0932486 0.0823065 -3.308445 -1.122715
## 6 322.65 478.67 286.91 306.70 0.0973228 0.0800806 -3.359507 -1.120668
## 7 334.54 477.58 292.09 311.91 0.0801688 0.0778710 -3.335309 -1.209756
## 8 210.75 381.72 247.38 213.17 0.0561224 0.0870046 -4.110956 -1.595047
##      lprbconv  lprbpris  lavgsen  lpolpc  ldensity  lwcon  lwtrd
## 1 -0.9111490 -0.7503061 1.724551 -6.327340 0.8360171 5.330205 5.810005
## 2 -0.8370060 -0.6792581 1.720979 -6.338704 0.8459773 5.360137 5.911600
## 3 -0.6430188 -0.7345839 1.757858 -6.300291 0.8509204 5.392628 7.240509
## 4 -0.5030129 -0.6537265 1.930071 -6.273361 0.8528909 5.409070 5.988612
## 5 -0.5469313 -0.6990466 1.879465 -6.253162 0.8607340 5.496169 5.882718
## 6 -0.6687981 -0.8212920 1.931521 -6.268420 0.8694848 5.552626 5.912277
## 7 -0.6394244 -0.8297232 1.903599 -6.304609 0.8848549 5.639869 6.013041
## 8 -0.1403569 -0.7640998 2.134166 -7.428766 -0.0234386 5.240520 5.678950
##      lwtrd  lwfir  lwser  lwmfgr  lwfed  lwsta  lwloc  lpctymle
## 1 5.205835 5.607452 5.374044 5.434246 6.014619 5.464848 5.444450 -2.433870
## 2 5.244607 5.706707 5.444911 5.482013 6.039540 5.536862 5.467174 -2.449038
```

```
## 3 5.281372 5.736475 5.481292 5.597310 6.084157 5.522900 5.515765 -2.464036
## 4 5.301128 5.858180 5.531204 5.640985 6.129421 5.568077 5.577387 -2.478925
## 5 5.332152 5.948220 5.564850 5.700042 6.195282 5.639919 5.664972 -2.497306
## 6 5.386862 6.015875 5.596987 5.776568 6.171011 5.659169 5.725870 -2.524721
## 7 5.399384 6.116272 5.613776 5.812757 6.168732 5.677062 5.742715 -2.552702
## 8 5.020080 5.310390 5.254230 5.350673 5.944687 5.510926 5.362090 -2.441794
##      lpctmin  ltaxpc      lmix
## 1 3.006608 3.246399 -2.303407
## 2 3.006608 3.213833 -2.272549
## 3 3.006608 3.275311 -2.517281
## 4 3.006608 3.289981 -2.544612
## 5 3.006608 3.337204 -2.372487
## 6 3.006608 3.392526 -2.329722
## 7 3.006608 3.433783 -2.523621
## 8 2.068926 2.678338 -2.880219
```

```
#Display all variable names
print(ls(Crime))
```

```
## [1] "avgsen" "county" "crmte" "density" "lavgsen" "lcrmte"
## [7] "ldensity" "lmix" "lpctmin" "lpctymle" "lpolpc" "lprbarr"
## [13] "lprbconv" "lprbpris" "ltaxpc" "lwcon" "lwfed" "lwfir"
## [19] "lwloc" "lwmfg" "lwser" "lwsta" "lwtrd" "lwtuc"
## [25] "mix" "pctmin" "pctymle" "polpc" "prbarr" "prbconv"
## [31] "prbpris" "region" "smsa" "taxpc" "wcon" "wfed"
## [37] "wfir" "wloc" "wmfg" "wser" "wsta" "wtrd"
## [43] "wtuc" "year"
```

```
#Display number of variables
print(length(ls(Crime)))
```

```
## [1] 44
```

```
#Display each variable's descriptive summary
print(summary(Crime))
```

```
##      county      year      crmte      prbarr
##  Min.   : 1.0    Min.   :81    Min.   :0.001812    Min.   :0.05882
##  1st Qu.: 51.0    1st Qu.:82    1st Qu.:0.018352    1st Qu.:0.21790
##  Median :103.0    Median :84    Median :0.028441    Median :0.27824
##  Mean   :100.6    Mean   :84    Mean   :0.031588    Mean   :0.30737
##  3rd Qu.:151.0    3rd Qu.:86    3rd Qu.:0.038406    3rd Qu.:0.35252
##  Max.   :197.0    Max.   :87    Max.   :0.163835    Max.   :2.75000
##      prbconv      prbpris      avgsen      polpc
##  Min.   : 0.06838    Min.   :0.1489    Min.   : 4.220    Min.   :0.0004585
##  1st Qu.: 0.34769    1st Qu.:0.3744    1st Qu.: 7.160    1st Qu.:0.0011913
##  Median : 0.47437    Median :0.4286    Median : 8.495    Median :0.0014506
##  Mean   : 0.68862    Mean   :0.4255    Mean   : 8.955    Mean   :0.0019168
##  3rd Qu.: 0.63560    3rd Qu.:0.4832    3rd Qu.:10.197    3rd Qu.:0.0018033
##  Max.   :37.00000    Max.   :0.6786    Max.   :25.830    Max.   :0.0355781
##      density      taxpc      region      smsa      pctmin
##  Min.   :0.1977    Min.   : 14.30    other :245    no :574    Min.   : 1.284
```

##	1st Qu.:0.5329	1st Qu.: 23.43	west :147	yes: 56	1st Qu.:10.005
##	Median :0.9526	Median : 27.79	central:238		Median :24.852
##	Mean :1.3861	Mean : 30.24			Mean :25.713
##	3rd Qu.:1.5078	3rd Qu.: 33.27			3rd Qu.:38.223
##	Max. :8.8277	Max. :119.76			Max. :64.348
##	wcon	wtuc	wtrd	wfir	
##	Min. : 65.62	Min. : 28.86	Min. : 16.87	Min. : 3.516	
##	1st Qu.: 201.66	1st Qu.: 317.60	1st Qu.: 168.05	1st Qu.:235.705	
##	Median : 236.46	Median : 358.20	Median : 185.48	Median :264.423	
##	Mean : 245.67	Mean : 406.10	Mean : 192.82	Mean :272.059	
##	3rd Qu.: 269.69	3rd Qu.: 411.02	3rd Qu.: 204.82	3rd Qu.:302.440	
##	Max. :2324.60	Max. :3041.96	Max. :2242.75	Max. :509.466	
##	wser	wmfg	wfed	wsta	
##	Min. : 1.844	Min. :101.8	Min. :255.4	Min. :173.0	
##	1st Qu.: 191.319	1st Qu.:234.0	1st Qu.:361.5	1st Qu.:258.2	
##	Median : 216.475	Median :271.6	Median :404.0	Median :289.4	
##	Mean : 224.671	Mean :285.2	Mean :403.9	Mean :296.9	
##	3rd Qu.: 247.155	3rd Qu.:320.0	3rd Qu.:444.6	3rd Qu.:331.5	
##	Max. :2177.068	Max. :646.9	Max. :598.0	Max. :548.0	
##	wloc	mix	pctymle	lcrmrte	
##	Min. :163.6	Min. :0.002457	Min. :0.06216	Min. : -6.314	
##	1st Qu.:226.8	1st Qu.:0.075324	1st Qu.:0.07859	1st Qu.: -3.998	
##	Median :253.1	Median :0.102089	Median :0.08316	Median : -3.560	
##	Mean :258.0	Mean :0.139396	Mean :0.08897	Mean : -3.609	
##	3rd Qu.:289.3	3rd Qu.:0.149009	3rd Qu.:0.08919	3rd Qu.: -3.260	
##	Max. :388.1	Max. :4.000000	Max. :0.27436	Max. : -1.809	
##	lprbarr	lprbconv	lprbpris	lavgsen	
##	Min. : -2.833	Min. : -2.6827	Min. : -1.9042	Min. :1.440	
##	1st Qu.: -1.524	1st Qu.: -1.0564	1st Qu.: -0.9824	1st Qu.:1.969	
##	Median : -1.279	Median : -0.7458	Median : -0.8473	Median :2.139	
##	Mean : -1.274	Mean : -0.6929	Mean : -0.8786	Mean :2.153	
##	3rd Qu.: -1.043	3rd Qu.: -0.4532	3rd Qu.: -0.7273	3rd Qu.:2.322	
##	Max. : 1.012	Max. : 3.6109	Max. : -0.3878	Max. :3.252	
##	lpolpc	ldensity	lwcon	lwtuc	
##	Min. : -7.688	Min. : -1.62091	Min. :4.184	Min. :3.362	
##	1st Qu.: -6.733	1st Qu.: -0.62934	1st Qu.:5.307	1st Qu.:5.761	
##	Median : -6.536	Median : -0.04857	Median :5.466	Median :5.881	
##	Mean : -6.491	Mean : -0.01593	Mean :5.463	Mean :5.916	
##	3rd Qu.: -6.318	3rd Qu.: 0.41066	3rd Qu.:5.597	3rd Qu.:6.019	
##	Max. : -3.336	Max. : 2.17789	Max. :7.751	Max. :8.020	
##	lwtrd	lwfir	lwser	lwmfg	
##	Min. :2.826	Min. :1.257	Min. :0.6118	Min. :4.623	
##	1st Qu.:5.124	1st Qu.:5.463	1st Qu.:5.2539	1st Qu.:5.455	
##	Median :5.223	Median :5.578	Median :5.3775	Median :5.604	
##	Mean :5.232	Mean :5.579	Mean :5.3646	Mean :5.615	
##	3rd Qu.:5.322	3rd Qu.:5.712	3rd Qu.:5.5100	3rd Qu.:5.768	
##	Max. :7.715	Max. :6.233	Max. :7.6857	Max. :6.472	
##	lwfed	lwsta	lwloc	lpctymle	
##	Min. :5.543	Min. :5.153	Min. :5.097	Min. : -2.778	
##	1st Qu.:5.890	1st Qu.:5.554	1st Qu.:5.424	1st Qu.: -2.543	
##	Median :6.001	Median :5.668	Median :5.534	Median : -2.487	
##	Mean :5.989	Mean :5.678	Mean :5.540	Mean : -2.443	
##	3rd Qu.:6.097	3rd Qu.:5.804	3rd Qu.:5.667	3rd Qu.: -2.417	
##	Max. :6.394	Max. :6.306	Max. :5.961	Max. : -1.293	

```
##      lpctmin      ltaxpc      lmix
## Min.   :0.2497   Min.    :2.660   Min.    :-6.009
## 1st Qu.:2.3030   1st Qu.:3.154   1st Qu.: -2.586
## Median :3.2127   Median :3.325   Median : -2.282
## Mean   :2.9134   Mean    :3.356   Mean    :-2.234
## 3rd Qu.:3.6434   3rd Qu.:3.505   3rd Qu.: -1.904
## Max.   :4.1643   Max.    :4.786   Max.    : 1.386
```

- b) [5 Pts] Calculate the mean, variance and standard deviation of probability of arrest (prbarr) by omitting the missing values, if any.

```
#mean
prbarr_avg <- mean(Crime$prbarr, na.rm = TRUE)
prbarr_avg
```

```
## [1] 0.3073682
```

```
#variance
prbarr_var <- var(Crime$prbarr, na.rm = TRUE)
prbarr_var
```

```
## [1] 0.02931104
```

```
#standard deviation
prbarr_sd <- sqrt(prbarr_var)
prbarr_sd
```

```
## [1] 0.1712047
```

- c) [5 Pts] Use lpolpc (log-police per capita) and smsa variables to build a linear regression model to predict probability of arrest (prbarr). And, compare with another linear regression model that uses polpc (police per capita) and smsa.

[5 Pts] How can you draw a conclusion from the results? (Note: Full marks requires comment on the predictors)

```
#Multiple linear regression
modell1 <- lm(prbarr ~ lpolpc + smsa, data = Crime)
summary(modell1)
```

```
##
## Call:
## lm(formula = prbarr ~ lpolpc + smsa, data = Crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.46050 -0.07973 -0.01784  0.05390  2.24094
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.88964    0.08152  10.913  < 2e-16 ***
```

```
## lpolpc      0.08784    0.01246    7.048 4.80e-12 ***
## smsayes     -0.13638    0.02305   -5.918 5.38e-09 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.1623 on 627 degrees of freedom
## Multiple R-squared:  0.104, Adjusted R-squared:  0.1012
## F-statistic: 36.4 on 2 and 627 DF, p-value: 1.109e-15
```

```
model2 <- lm(prbarr ~ polpc + smsa, data = Crime)
summary(model2)
```

```
##
## Call:
## lm(formula = prbarr ~ polpc + smsa, data = Crime)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.72651 -0.07840 -0.01759  0.04955  2.22692
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.28213     0.00807  34.958 < 2e-16 ***
## polpc       18.34603     2.34684   7.817 2.29e-14 ***
## smsayes     -0.11163     0.02254  -4.953 9.40e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.161 on 627 degrees of freedom
## Multiple R-squared:  0.1189, Adjusted R-squared:  0.1161
## F-statistic: 42.31 on 2 and 627 DF, p-value: < 2.2e-16
```

#The p value of the F statistic in model1 and model2 are both highly significant. Model1 residuals have

*#Looking at the Coefficients from Model1, we can conclude that there is a highly significant (Pr(>|t|***

*#Looking at the Coefficients from Model2, we can conclude that there is a highly significant (Pr(>|t|***

- d) [5 Pts] Based on the output of your model, write the equations using the intercept and factors of smsa when polpc is set to 0.0015. and compare the result with predict() function.
Hint: Explore predict() function

```
#Model 2 equation when smsa is no
y.smsano = 0.28213 + 18.34603*(0.0015)
y.smsano
```

```
## [1] 0.309649
```

```
#Model2 equation when smsa is yes
y.smsayes = 0.28213 + 18.34603*(0.0015) - 0.11163
y.smsayes
```



```
## [1] 0.198019
```

```
#Temporary dataframe to set polpc to 0.0015
polpc.smsa.df <- data.frame(polpc = c(0.0015), smsa = c('no','yes'))
predict(model2, polpc.smsa.df)
```

```
##          1          2
## 0.3096441 0.1980168
```

#The results obtained from using the predict function are exactly the same up to 5 decimal points.

- e) [5 Pts] Find Pearson correlation between probability of prison sentence `prbpris` and tax per capita `taxpc`; and also Pearson correlation between probability of conviction `prbconv` and probability of arrest `prbarr`.

[5 Pts] What conclusions can you draw? Write your reasons as comments.

```
cor.test(Crime$prbpris, Crime$taxpc)
```

```
##
## Pearson's product-moment correlation
##
## data: Crime$prbpris and Crime$taxpc
## t = -2.8261, df = 628, p-value = 0.004862
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## -0.18852675 -0.03424894
## sample estimates:
##          cor
## -0.1120631
```

#We can conclude that there is a very low negative linear correlation between prison sentence (prbpris)

```
cor.test(Crime$prbconv, Crime$prbarr)
```

```
##
## Pearson's product-moment correlation
##
## data: Crime$prbconv and Crime$prbarr
## t = 0.89192, df = 628, p-value = 0.3728
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## -0.04266359 0.11336788
## sample estimates:
##          cor
## 0.0355689
```

#We can conclude that there is a very low positive correlation linear relationship between the probabil

- f) [5 Pts] Display the correlation matrix of the variables: `prbconv`, `prbpris`, `avgsen`, `polpc`.

[5 Pts] Write what conclusion you can draw, as comments.

```
#install.packages("corrplot")
library(corrplot)
```

```
## Warning: package 'corrplot' was built under R version 4.3.2
```

```
## corrplot 0.92 loaded
```

```
table_cor<- cor(Crime[,5:8])
table_cor
```

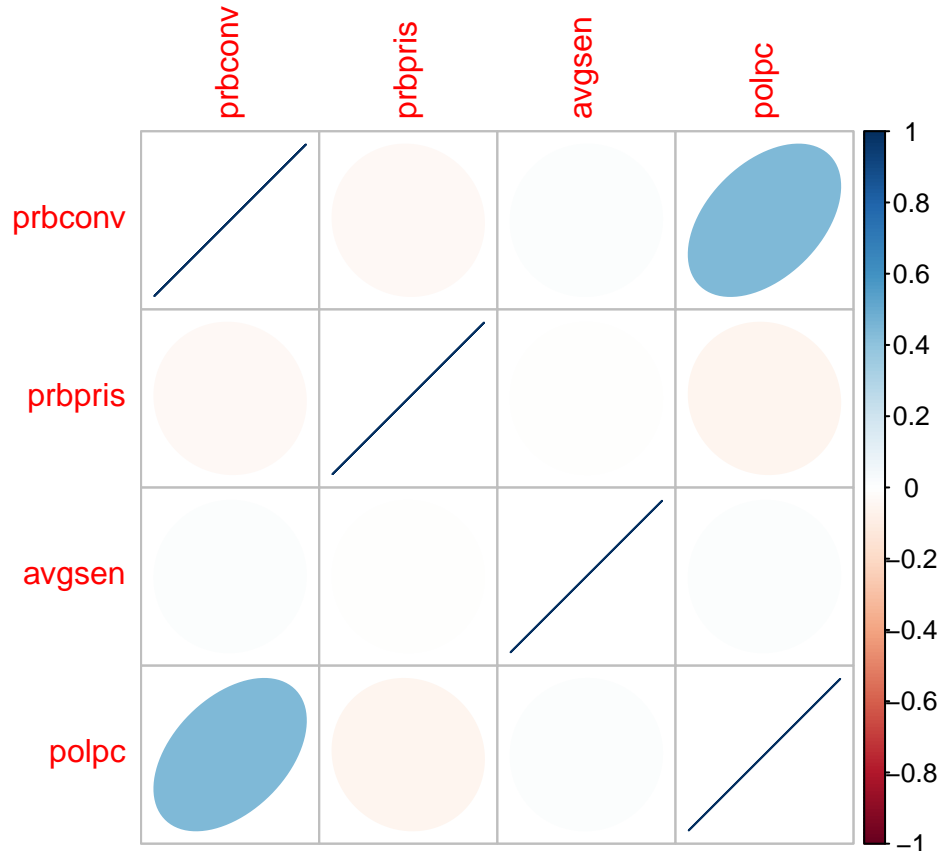
```
##           prbconv      prbpris      avgsen      polpc
## prbconv  1.00000000 -0.037340175  0.015304708  0.44963500
## prbpris -0.03734017  1.000000000 -0.004299394 -0.05745238
## avgsen   0.01530471 -0.004299394  1.000000000  0.01712970
## polpc    0.44963500 -0.057452385  0.017129699  1.00000000
```

#We can conclude that there is a moderate positive correlation coefficient between prbconv and polpc (0.45)

#We can also conclude that there is very little negative linear relationship between prbconv and prbpris (-0.037)

#We can also conclude that there is a very low positive correlation between prbconv and avgsen (0.015)

```
#Further visualizing pearson correlations
corrplot(table_cor, method = "ellipse")
```



Question 3 [15 Pts]

This question makes use of package “ISwR”. Please load `airquality` dataset as following:

```
#install.packages("ISwR")  
library(ISwR)
```


```
## Warning: package 'ISwR' was built under R version 4.3.2
```

```
data(airquality)  
str(airquality)
```

```
## 'data.frame': 153 obs. of 6 variables:  
## $ Ozone : int 41 36 12 18 NA 28 23 19 8 NA ...  
## $ Solar.R: int 190 118 149 313 NA NA 299 99 19 194 ...  
## $ Wind : num 7.4 8 12.6 11.5 14.3 14.9 8.6 13.8 20.1 8.6 ...  
## $ Temp : int 67 72 74 62 56 66 65 59 61 69 ...  
## $ Month : int 5 5 5 5 5 5 5 5 5 5 ...  
## $ Day : int 1 2 3 4 5 6 7 8 9 10 ...
```

- a) [5 Pts] Plot a histogram to assess the normality of the `Solar.R` variable, then explain why it does not appear normally distributed.

```
print(hist(airquality$Solar.R))
```



```
## $breaks
## [1] 0 50 100 150 200 250 300 350
##
## $counts
## [1] 17 17 18 19 30 36 9
##
## $density
## [1] 0.002328767 0.002328767 0.002465753 0.002602740 0.004109589 0.004931507
## [7] 0.001232877
##
## $mids
## [1] 25 75 125 175 225 275 325
##
## $xname
## [1] "airquality$Solar.R"
##
## $equidist
## [1] TRUE
##
## attr("class")
## [1] "histogram"
```

#This variable does not appear normally distributed because the histogram is not a bell shaped curve (t


- b) [5 Pts] Create a boxplot that shows the distribution of `Solar.R` in each month. Use different colors for each month.

```
#Colour vector preparation
library(RColorBrewer)
```

```
boxplot(Solar.R ~ Month, data = airquality, main = "Distribution of Solar.R by Month", xlab = "Months (1
```


```
## Warning in brewer.pal(12, name = "Set2"): n too large, allowed maximum for palette Set2 is 8
## Returning the palette you asked for with that many colors
```

Distribution of Solar.R by Month



- c) [5 Pts] Create a matrix of scatterplots of all the numeric variables in the `airquality` dataset (i.e. Ozone, Solar.R, Wind and Temp.) (Hint: investigate `pairs()` function)

```
#Matrix only comparing: Ozone, Solar.R, Wind and Temp
print(pairs(~ Ozone + Solar.R + Wind + Temp, data = airquality, main = "Airquality Data", na.action = na.omit))
```



```
## NULL
```

Question 4 [25 Pts]

Many times in data analysis, we need a method that relies on repeated random sampling to obtain numerical results. The underlying concept is to use randomness to solve problems. In fact, this is a mathematical technique, which is used to estimate the possible outcomes of an uncertain event and is called the *Monte Carlo Method*.

Consider that We roll a die 10 times and we want to know the probability of getting more than 3 times of even numbers. This is a problem for the Binomial distribution, but suppose we don't know anything about Binomial distribution. We can easily solve this problem with a Monte Carlo Simulation.

- a) [5 Pts] The Monte Carlo Method uses random numbers to simulate some process. Here the process is rolling a die 10 times. Assume the die is fair. What is the probability of success or getting an even number in rolling the die once?

```
#install.packages("gtools")
library(gtools)
```

```
## Warning: package 'gtools' was built under R version 4.3.2
```

```

one.dice <- c(1, 2, 3, 4, 5, 6)
perm <- permutations(length(one.dice), 1, one.dice,
repeats.allowed = TRUE)
perm

```

```

##      [,1]
## [1,]    1
## [2,]    2
## [3,]    3
## [4,]    4
## [5,]    5
## [6,]    6

```

#Probability of success of rolling an even number when rolling the die once is 3/6 or 0.5

- b) [10 Pts] Define a function named `one.trial`, that simulates a single round of rolling a die 10 times and returns true if the number of even numbers is > 3 .

#INSERT YOUR ANSWER HERE.

```

one.trial <- function(){

  die <- c()
  num_even <- 0

  #Roll the die and store rolls in a vector called "die"
  for (roll in 1:10){
    die <- append(die, sample(1:6, size = 1, replace = TRUE))
  }
  #print(die)

  #Count number of even numbers in the vector "die"
  for (i in die) {
    if (i == 2) {
      num_even <- num_even + 1
    } else if (i == 4){
      num_even <- num_even + 1
    } else if (i == 6){
      num_even <- num_even + 1
    } else {
      num_even <- num_even + 0
    }
  }

  #print(num_even)

  #Determine whether count of even numbers is over 3
  if (num_even > 3) {
    return(TRUE)
  } else{
    return(FALSE)
  }

}

```

```
one.trial()
```

```
## [1] TRUE
```

- c) [5 pts] Repeat the function `one.trial` for $N = 100,000$ times and sum up the outcomes and store the result in a variable named `desired.output`. Compute the probability of getting more than 3 times of even numbers by using relative frequency.

```
#Returns the number one.trial() that is equal to TRUE (>3 even numbers)
set.seed(10)
desired.output <- sum(replicate(n = 100000, expr = one.trial()))
desired.output
```

```
## [1] 82924
```

```
my.probability <- desired.output/100000
my.probability
```

```
## [1] 0.82924
```

- d) [5 pts] Use the Binomial formula you learned before to calculate such probability and Compare it with the probability value obtained in part (c).

```
set.seed(10)
pbinom(q = 3, size = 10, prob = 1/2)
```

```
## [1] 0.171875
```

Congratulations! you have completed the first run of the Monte Carlo simulation.

If there is further interest, put all the above logic in a function, and call it 50 times at least, and store the results in a vector called `Prob` then take the mean of `Prob` vector to be more accurate.

**** End of Assignment ****